



GPT-5.6 latency scaling is robust across estimators

Quadratic-capable fits using three treatments of request-level latency noise



• Raw request — Student-t - - Frontier - - Spike + contention