

Illustrative latency extrapolation to 10 million tokens

Quadratic Student-t fits for GPT-5.6; linear Student-t fits for Claude 5

